

# Frederick National Laboratory for Cancer Research

sponsored by the National Cancer Institute



## Generative Adversarial Networks

Dr. Yanling Liu  
Manager, Imaging and Visualization Group, ABCS, FNLCR

DEPARTMENT OF HEALTH AND HUMAN SERVICES • National Institutes of Health • National Cancer Institute

Frederick National Laboratory is a Federally Funded Research and Development Center operated by Leidos Biomedical Research, Inc., for the National Cancer Institute

**FNLCR**

## **Imaging and Visualization Group (IVG)**

---

***We convert images to knowledge for cancer and biomedical research***

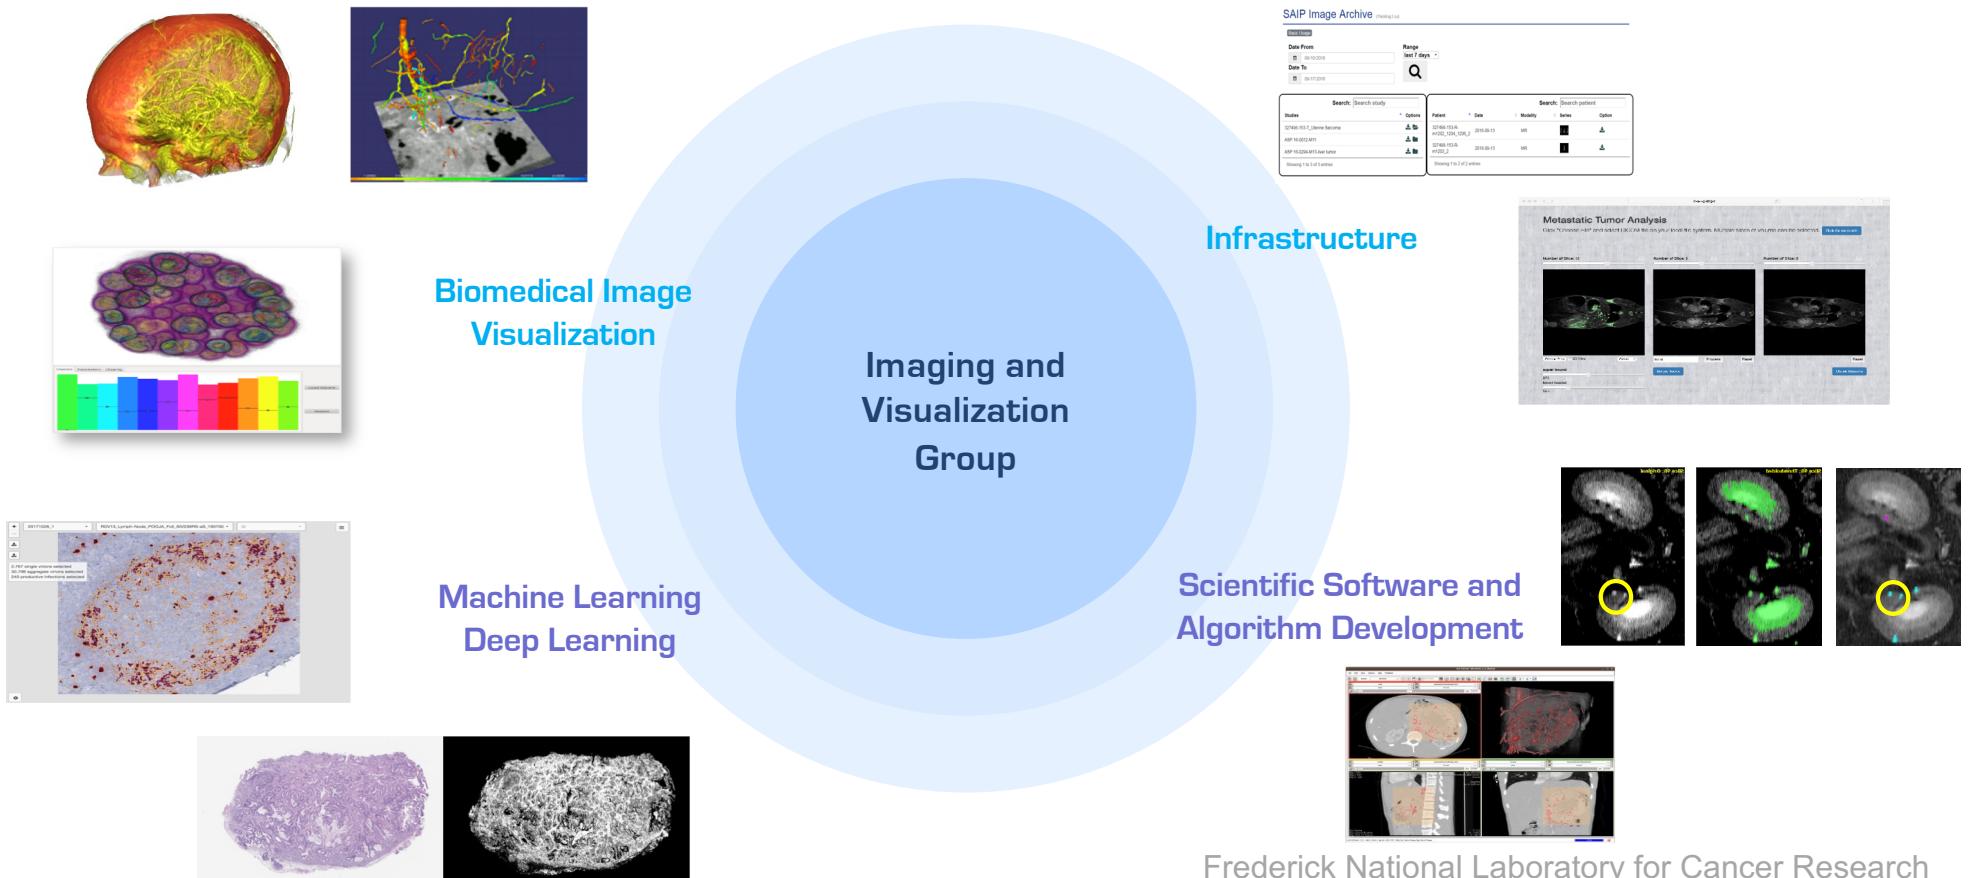
***Building systems that automate image processing***

***Improving efficiency and enhancing consistency***

***Creating better tools and workflows for researchers***

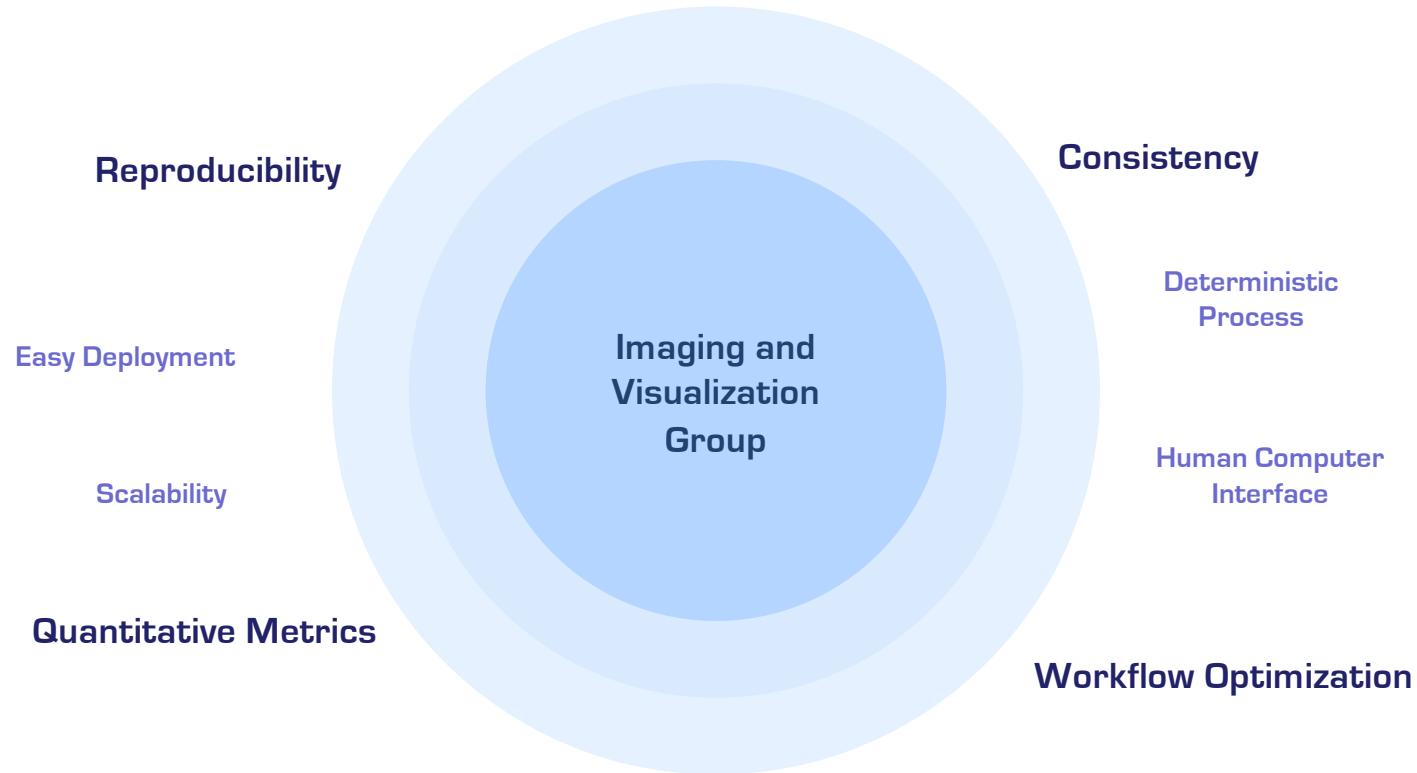
Frederick National Laboratory for Cancer Research

# Frederick National Laboratory for Cancer Research Imaging and Visualization Group (IVG)



# Frederick National Laboratory for Cancer Research

## Imaging and Visualization Group (IVG)



***“...however, image interpretation by humans is limited due to its subjectivity, large variations across interpreters, and fatigue.”***

*H. Greenspan, B. van Ginneken and R. M. Summers, "Guest Editorial Deep Learning in Medical Imaging: Overview and Future Promise of an Exciting New Technique," in IEEE Transactions on Medical Imaging, vol. 35, no. 5, pp. 1153-1159, May 2016.*

---

To enhance image **interpretation** by **humans**  
*via **quantitative** information derived from **Deep Learning** based image **segmentation***

Frederick National Laboratory for Cancer Research

# Machine Learning and Deep Learning

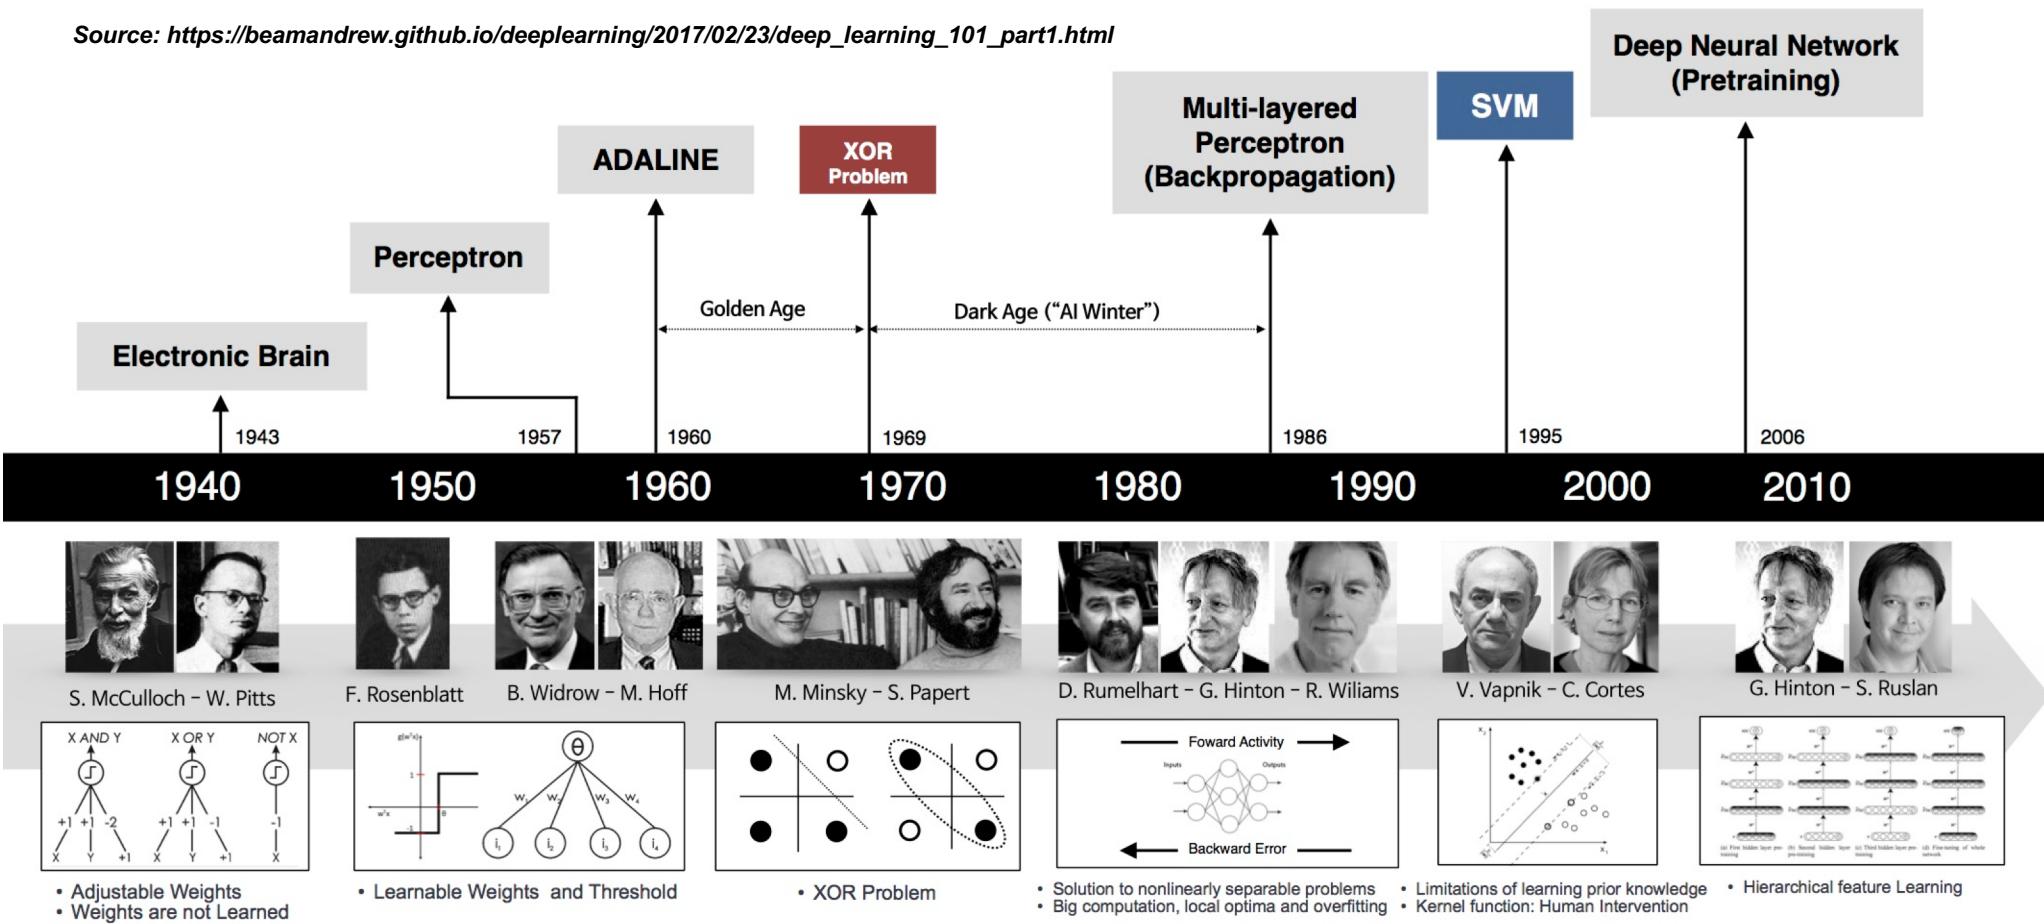
---

*A little bit of history...*

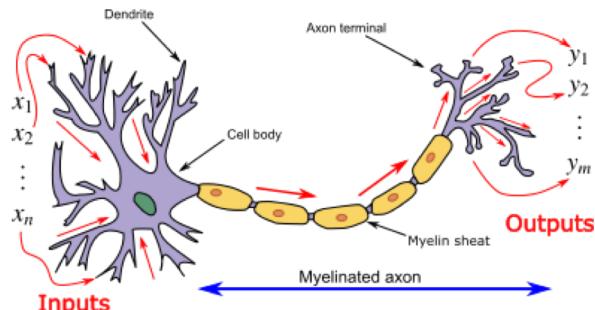
Frederick National Laboratory for Cancer Research

# Machine Learning: A Brief History

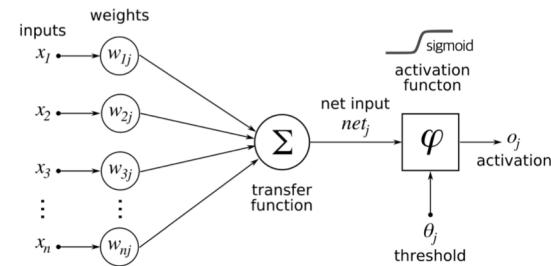
Source: [https://beamandrew.github.io/deeplearning/2017/02/23/deep\\_learning\\_101\\_part1.html](https://beamandrew.github.io/deeplearning/2017/02/23/deep_learning_101_part1.html)



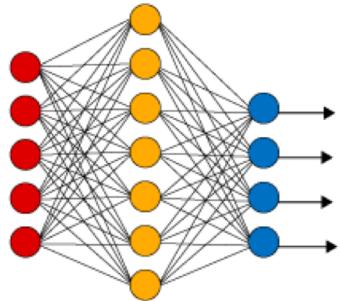
# A Simplified View of Deep Learning



*Image credit: Wikipedia*



Simple Neural Network

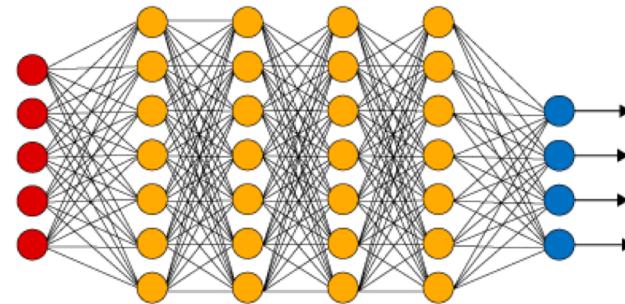


● Input Layer

○ Hidden Layer

● Output Layer

Deep Learning Neural Network



*Image credit: Deep Learning made easy with Deep Cognition from Medium*

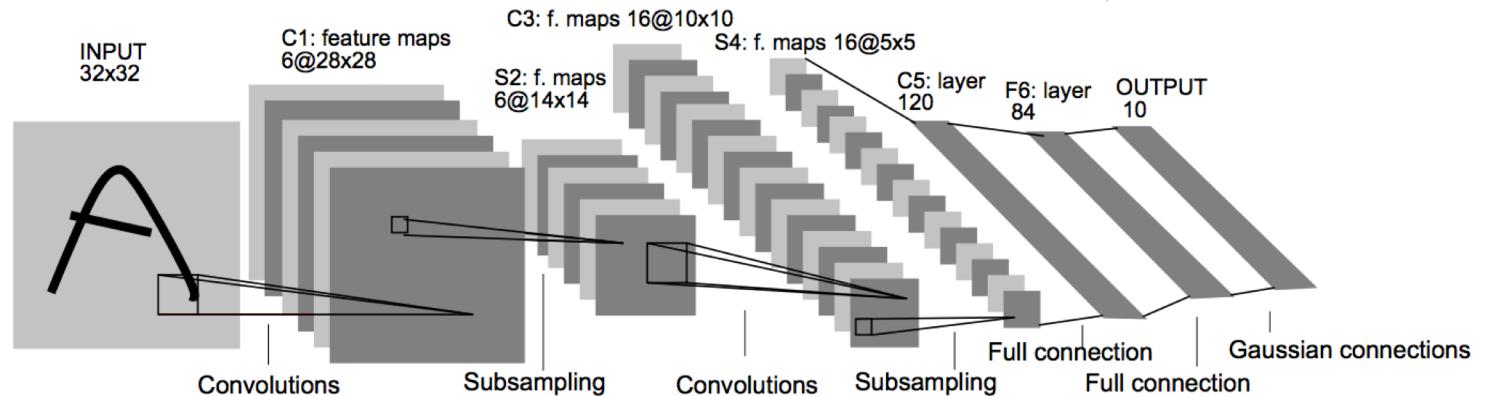
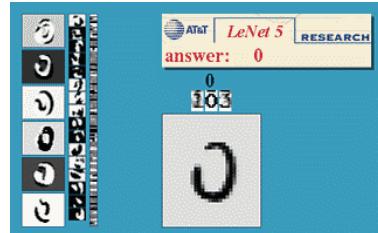
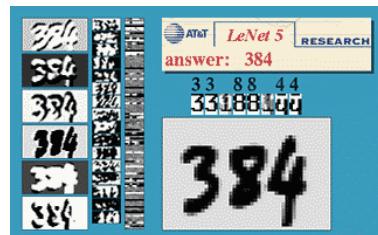
# Example of Convolutional Neural Network: LeNet

Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner.

**Gradient-based learning applied to document recognition.**

*Proceedings of the IEEE*, November 1998.

Source: [http://machinelearningguru.com/computer\\_vision/basics/convolution/image\\_convolution\\_1.html](http://machinelearningguru.com/computer_vision/basics/convolution/image_convolution_1.html)



*LeNet 5 Architecture*

**Fun facts:** 2-3 days of training time on a Silicon Graphics Origin 2000 using a single 200MHz R10000 processor

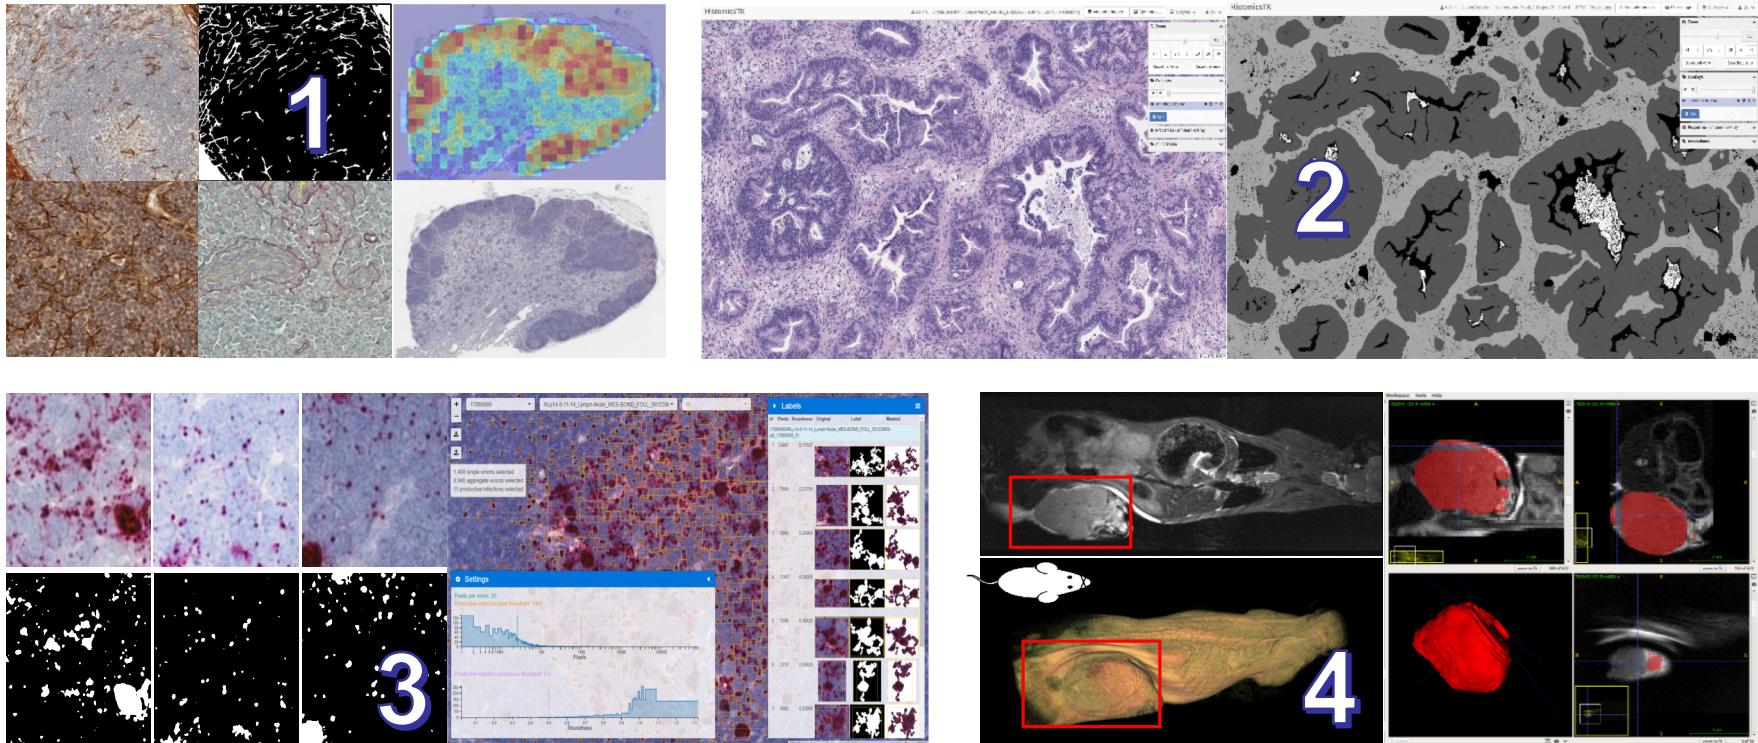
Frederick National Laboratory for Cancer Research

## The Training Data Challenge

---

Frederick National Laboratory for Cancer Research

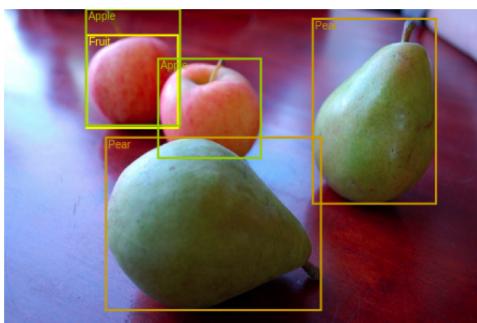
## Examples of Deep Learning enabled projects within IVG



# The Training Data Challenge



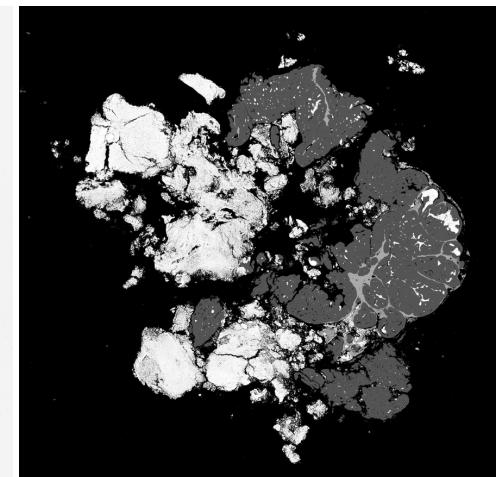
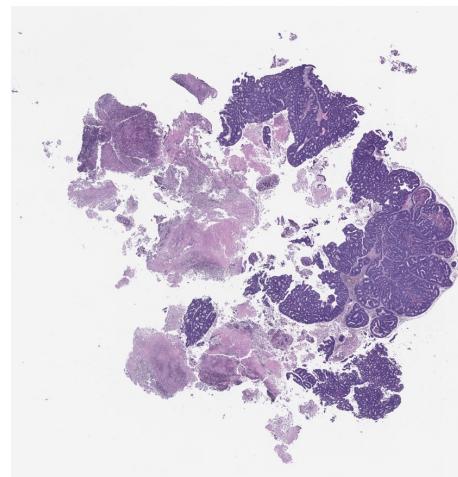
Source: <https://www.pyimagesearch.com/2018/05/07/multi-label-classification-with-keras/>



## Object Detection

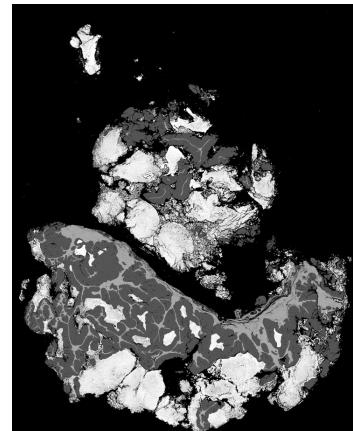
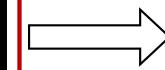
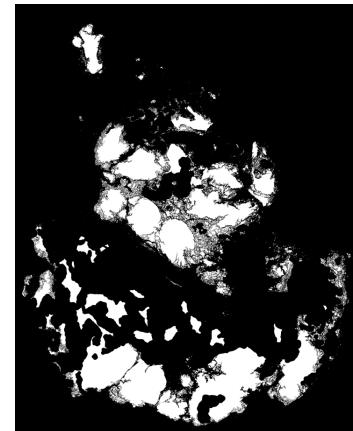
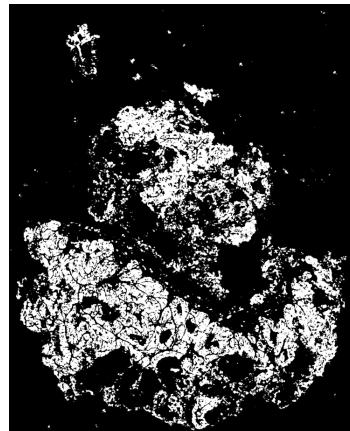
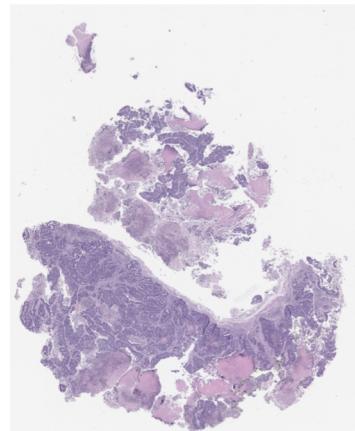
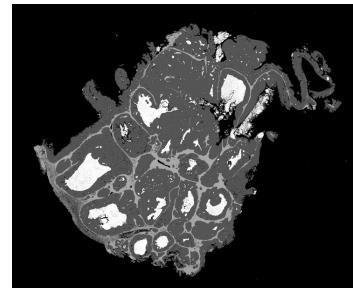
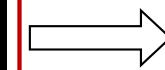
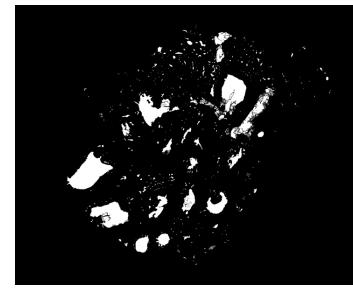
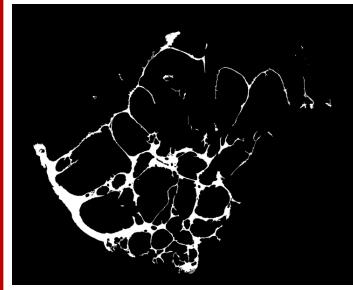
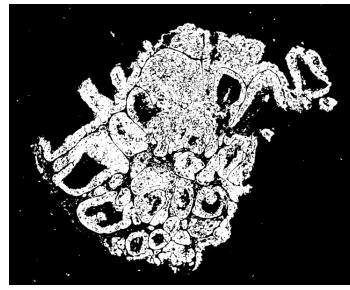
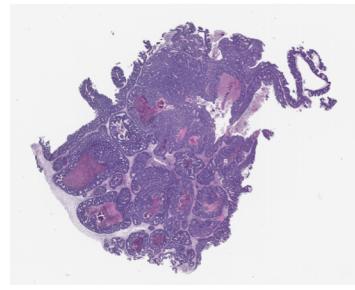
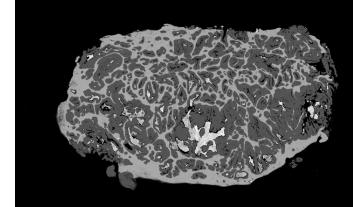
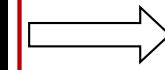
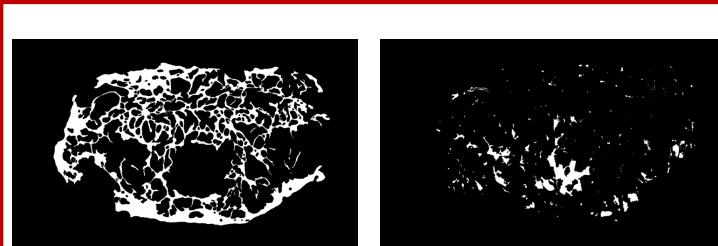
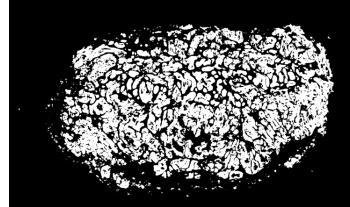
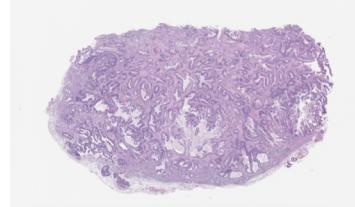
Source: <https://blog.algorithmia.com/deep-dive-into-object-detection-with-open-images-using-tensorflow/>

## Classification



## Precise Segmentation

*Collaboration with the MoCha group*



Three sample WSIs from  
the MoCha group

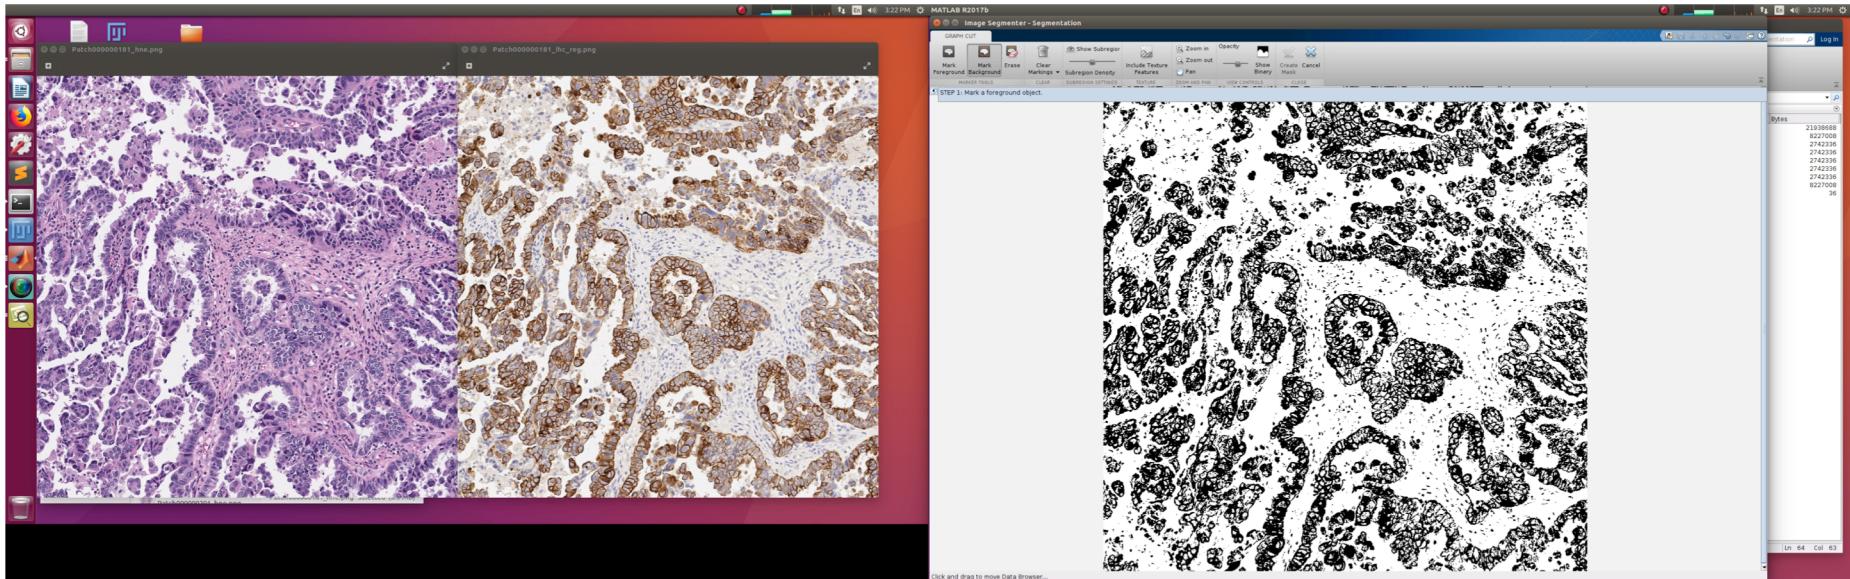
Epithelial cell annotations  
from the trained U-Net

Manual stroma annotation  
Guided by a pathologist

Combined annotations

# The Training Data Challenge

*Video showing tedious semi-auto pixel annotation creation process*



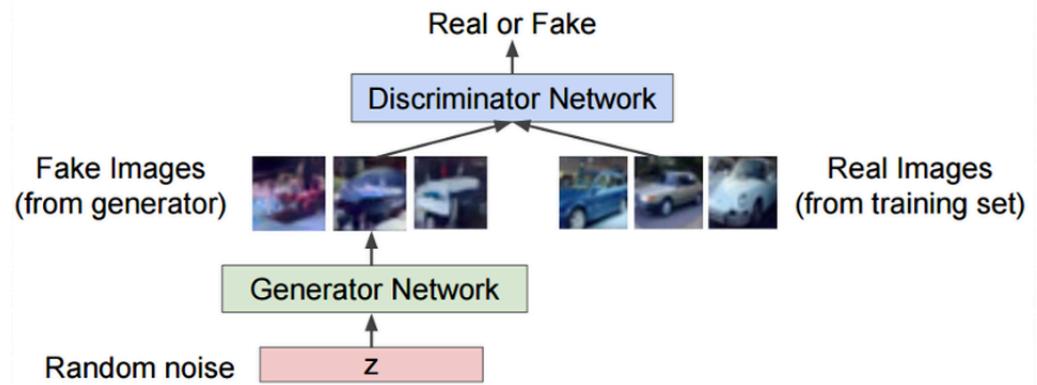
# **Generative Adversarial Networks (GANs)**

---

Frederick National Laboratory for Cancer Research

# GANs

- A system of two networks
  - Generator
  - Discriminator
  - Double feedback loops
    - The discriminator is in a feedback loop with the ground truth of the images, which we know.
    - The generator is in a feedback loop with the discriminator.



GANs in action. (Source - <https://goo.gl/NAN1Jw>)



## Applications of GANs



Source



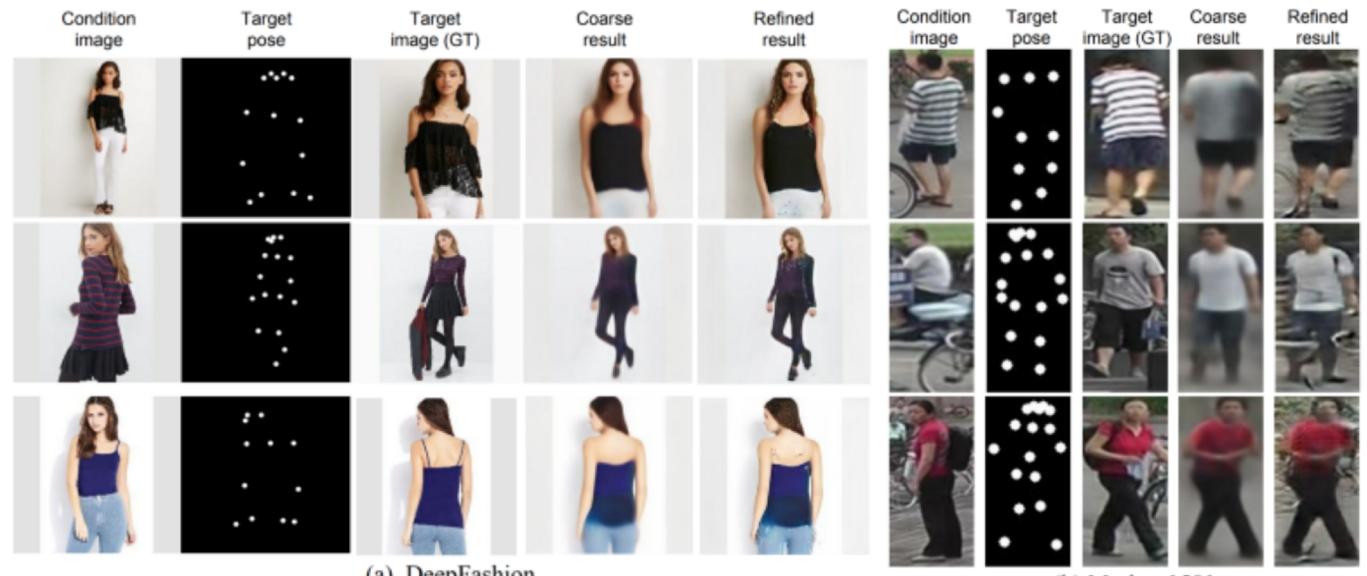
Figure 7: Generated samples

Towards the automatic Anime characters creation with Generative Adversarial Networks

# Applications of GANs

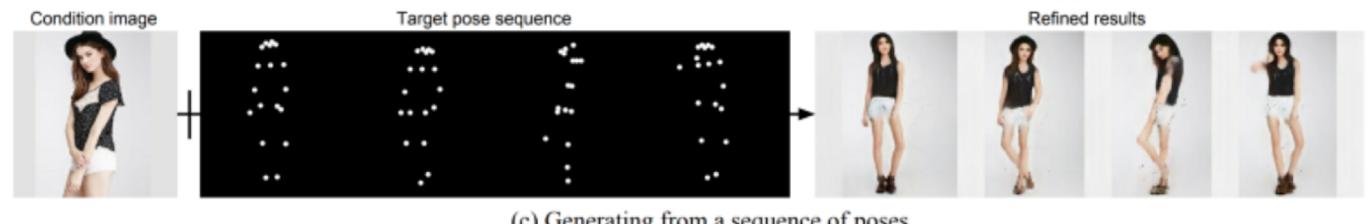


Source



(a) DeepFashion

(b) Market-1501



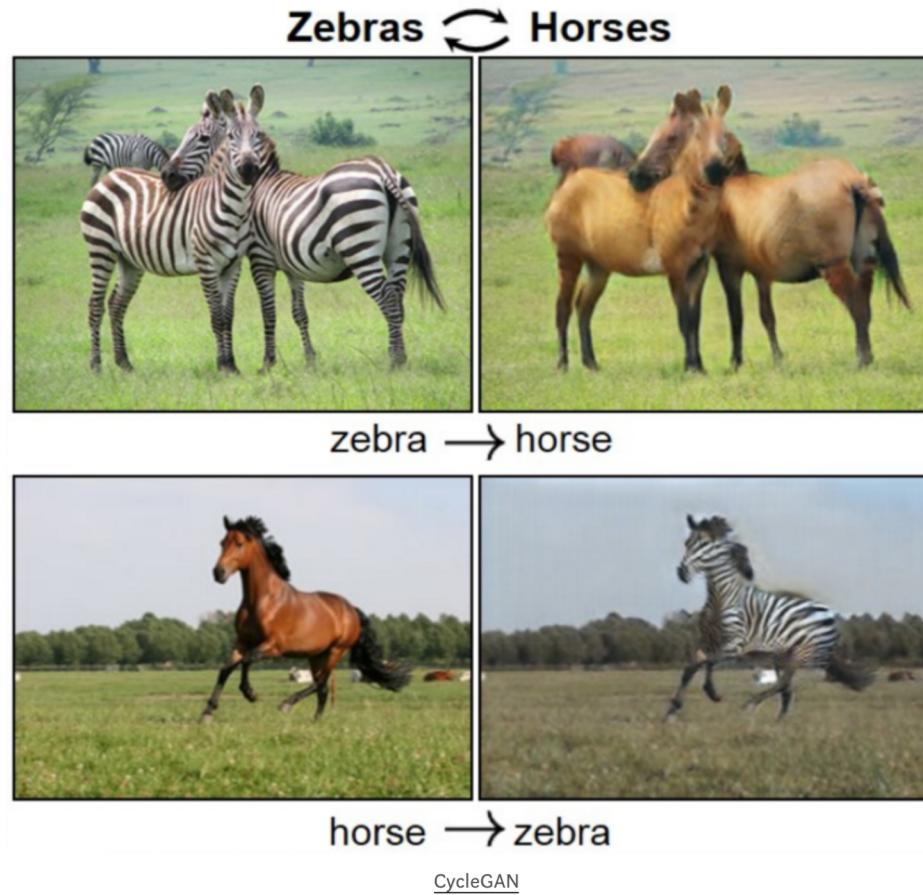
(c) Generating from a sequence of poses

Pose Guided Person Image Generation

# Applications of GANs



Source



## Applications of GANs



Source



Figure 5:  $1024 \times 1024$  images generated using the CELEBA-HQ dataset. See Appendix F for a larger set of results, and the accompanying video for latent space interpolations.

[Progressive growing of GANs](#)

# Applications of GANs



Source

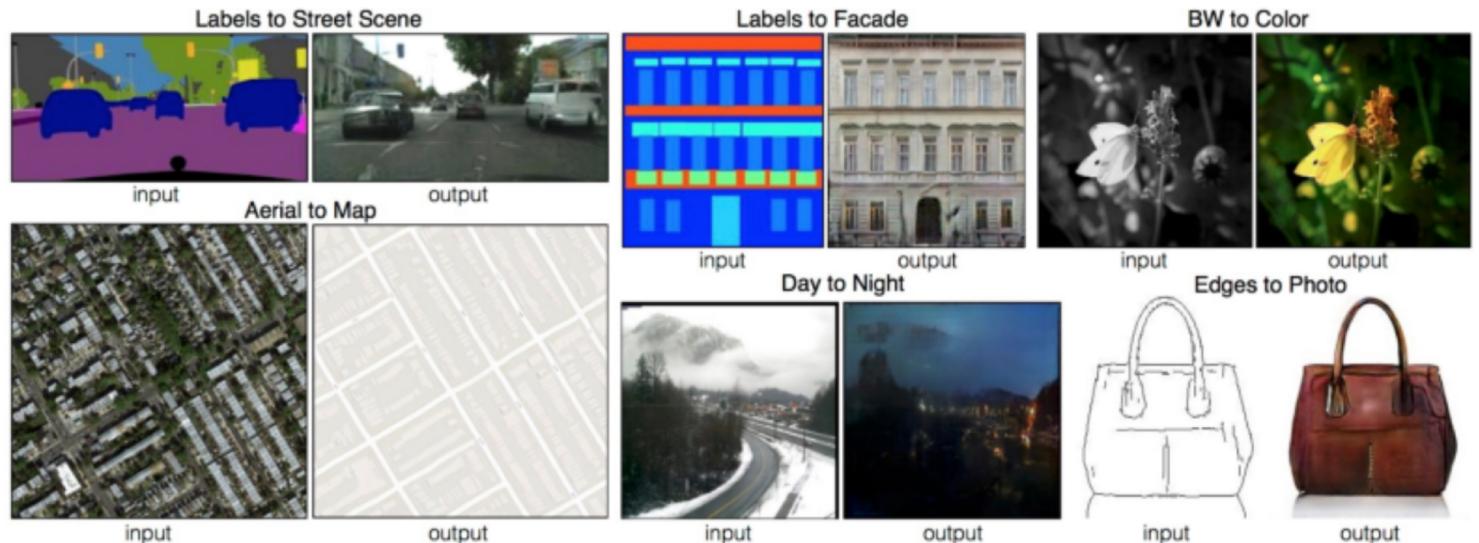


IcGAN

# Applications of GANs



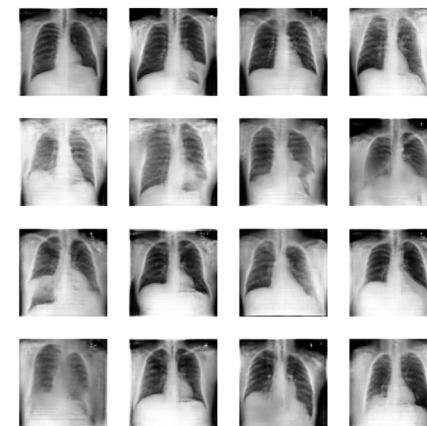
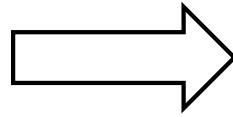
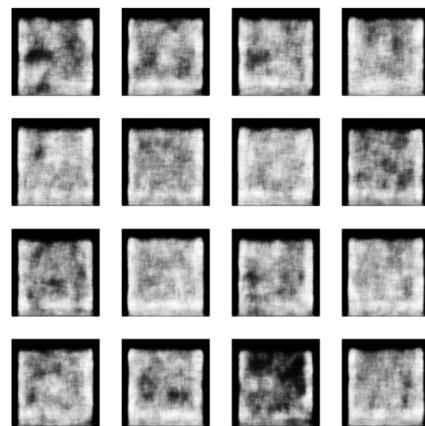
Source



pix2pix

## Tutorial: GAN for synthetic X-ray images

---



atory for Cancer Research

# RaGAN: modularizing GAN networks for easy data argumentation

- Radiology GAN (RaGAN)
  - GAN: Generative Adversarial Network
  - NIH Hackathon 2018, Sept. 10-12
    - <https://github.com/NCBI-Hackathons/RaGAN>
  - Attempt to modularize GAN networks for easy and quick (radiology) data argumentation
    - TensorFlow implementation
    - Wrapped as a Docker container
      - <https://hub.docker.com/r/johnbamboobilly/ragan/>
    - Default 6-layer minuscule network model (proof-of-concept)
    - Future plan
      - Allow drop-in network model integration for quick adoption

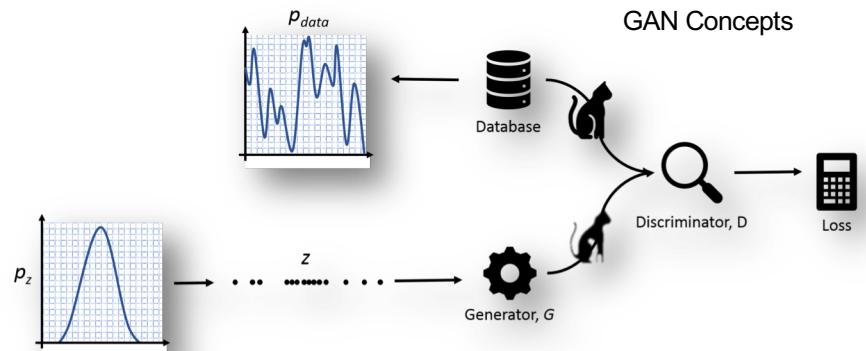
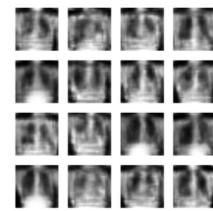


Image Source: <https://mlnotebook.github.io/img/CNN/gan1.png>



A chest X-ray image  
NIH CXR8



Synthesized images of the chest  
X-Ray data set, 2500 iterations, 28  
x 28 pixels each image, default  
minuscule network

## NIH CXR8 Dataset

- Publicly available since 09/2017
  - <https://nihcc.app.box.com/v/ChestXray-NIHCC>
  - More than 30,000 patients
  - Over 100,000 anonymized chest x-ray images

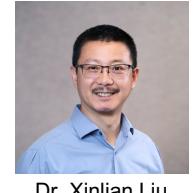
Nodule	2,705
Effusion	3,955
Atelectasis	4,215
Infiltration	9,547
No Finding	60,361
...	...

Examples of “outliers” in the dataset



# Data Preprocessing

- Source code
  - <https://github.com/blazers/gan/blob/master/scripts/p10book.ipynb>
  - [https://github.com/blazers/gan/blob/master/lists/cxr14\\_bad\\_labels.csv](https://github.com/blazers/gan/blob/master/lists/cxr14_bad_labels.csv)
  - [https://github.com/abcsFrederick/gan/blob/master/lists/Data\\_Entry\\_2017\\_parsed.csv](https://github.com/abcsFrederick/gan/blob/master/lists/Data_Entry_2017_parsed.csv)
- Group selection



Dr. Xinlian Liu

```
t_at=table[(table.label=='Atelectasis') & (table.view=='PA') & (table.gender=='M') & (table.age < 71) & (table.age > 19)]
```

```
t_in=table[(table.label=='Infiltration') & (table.view=='PA') & (table.gender=='M') & (table.age < 71) & (table.age > 19)]
```

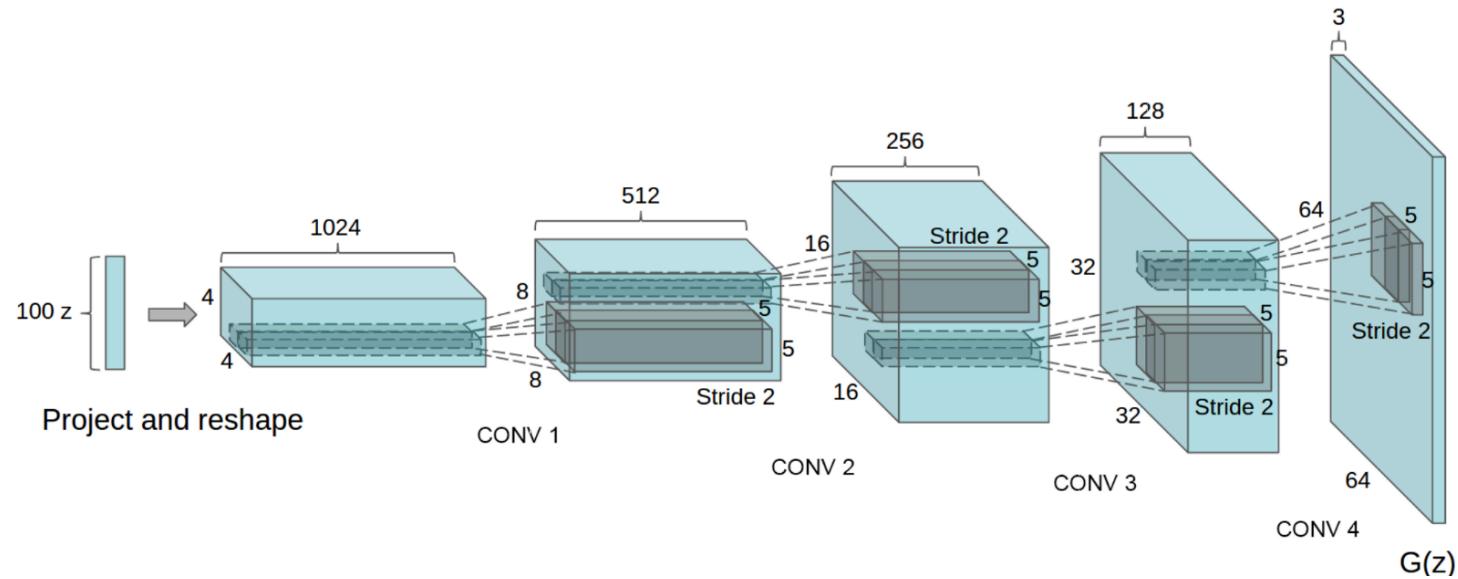
```
t_no=table[(table.label=='No Finding') & (table.view=='PA') & (table.gender=='M') & (table.age < 71) & (table.age > 19)]
```

```
: t_ef=table[(table.label=='Effusion') & (table.view=='PA') & (table.gender=='M') & (table.age < 71) & (table.age > 19)]
```

```
: t_nod=table[(table.label=='Nodule') & (table.view=='PA') & (table.gender=='M') & (table.age < 71) & (table.age > 19)]
```

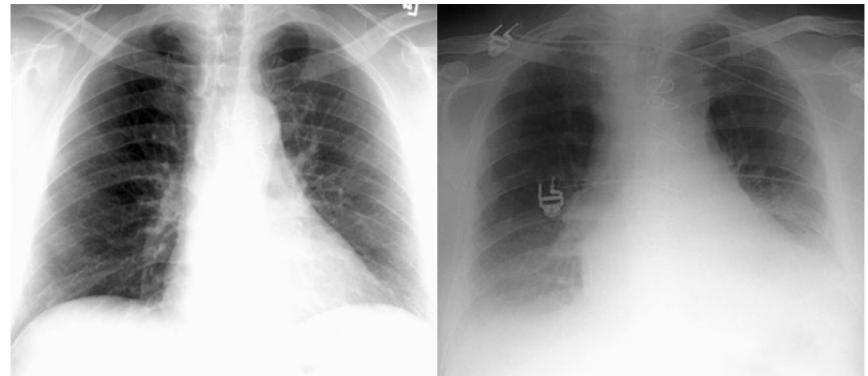
# DCGAN

- Original paper
  - <https://arxiv.org/abs/1511.06434>
  - Deep Convolutional Generative Adversarial Network (DCGAN)



# Test 1: GAN -> 64x64 synthetic images

- Training data
    - Infiltration subset
      - Anteriorposterior (AP) view
      - Male
      - Ages between 19 ~ 71
    - 2,126 samples
      - Down-sampled to 64x64 using imagemagick
  - Training statistics
    - Four Nvidia Quadro M6000 GPUs
    - 2,000 steps, 256 samples / step
    - Finished in 18 minutes



Posterioranterior (PA)                    Anteriorposterior (AP)  
<https://www.med-ed.virginia.edu/courses/rad/cxr/technique3chest.html>

## Source Code

- [https://github.com/abcsFrederick/gan/blob/master/scripts/NIH.AI\\_DCGAN\\_64x64\\_CXR8.py](https://github.com/abcsFrederick/gan/blob/master/scripts/NIH.AI_DCGAN_64x64_CXR8.py)

```
import os
import glob
import numpy as np
import time

from skimage.io import imread

import tensorflow as tf
from tensorflow.python.keras.models import Sequential
from tensorflow.python.keras.layers import Dense, Activation, Flatten, Reshape
from tensorflow.python.keras.layers import Conv2D, Conv2DTranspose, UpSampling2D
from tensorflow.python.keras.layers import LeakyReLU, Dropout
from tensorflow.python.keras.layers import BatchNormalization
from tensorflow.python.keras.optimizers import Adam, RMSprop

import matplotlib.pyplot as plt
```

## Source Code

- [https://github.com/abcsFrederick/gan/blob/master/scripts/NIH.AI\\_DCGAN\\_64x64\\_CXR8.py](https://github.com/abcsFrederick/gan/blob/master/scripts/NIH.AI_DCGAN_64x64_CXR8.py)

```
def load_data(directory_path):
    train_data = []
    for file_ in sorted(os.listdir(directory_path)):
        if file_.endswith('.png'):
            if directory_path.endswith('/'):
                image_path = directory_path + file_
            else: image_path = directory_path + '/' + file_

            #image = imread(image_path)/255.0 # Normalize values
            image = imread(image_path, as_gray=True)
            if np.max(image) > 1.:
                image = image / 255.0 #normalize values
            #print(image.shape)
            image = np.expand_dims(image, axis=-1) # Add channel dim
            train_data.append(image)

    train_data = np.array(train_data)
    return train_data
```

## Source

- <https://>

R8.py

```
# (W-F+2P)/S+1
def discriminator(self):
    if self.D:
        return self.D
    self.D = Sequential()
    depth = 64
    dropout = 0.2

    input_shape = (self.img_rows, self.img_cols, self.channel)
    self.D.add(Conv2D(depth*1, 5, strides=2, input_shape=input_shape, padding='same'))
    self.D.add(LeakyReLU(alpha=0.2))
    self.D.add(Dropout(dropout))

    self.D.add(Conv2D(depth*2, 5, strides=2, padding='same'))
    self.D.add(LeakyReLU(alpha=0.2))
    self.D.add(Dropout(dropout))

    self.D.add(Conv2D(depth*4, 5, strides=2, padding='same'))
    self.D.add(LeakyReLU(alpha=0.2))
    self.D.add(Dropout(dropout))

    self.D.add(Conv2D(depth*8, 5, strides=1, padding='same'))
    self.D.add(LeakyReLU(alpha=0.2))
    self.D.add(Dropout(dropout))

    # Out: 1-dim probability
    self.D.add(Flatten())
    self.D.add(Dense(1))
    self.D.add(Activation('sigmoid'))
    self.D.summary()
    return self.D
```

## Source Code

- [https://github.com/keras-team/keras/blob/master/examples/gan/gan\\_cifar10.py](https://github.com/keras-team/keras/blob/master/examples/gan/gan_cifar10.py)

[CXR8.py](#)

```
def generator(self):
    if self.G:
        return self.G
    self.G = Sequential()
    dropout = 0.2
    depth = 64+64+64+64
    dim = 8

    self.G.add(Dense(dim*dim*depth, input_dim=100))
    self.G.add(BatchNormalization(momentum=0.9))
    self.G.add(Activation('relu'))
    self.G.add(Reshape((dim, dim, depth)))
    self.G.add(Dropout(dropout))

    self.G.add(UpSampling2D())
    self.G.add(Conv2DTranspose(int(depth/2), 5, padding='same'))
    self.G.add(BatchNormalization(momentum=0.9))
    self.G.add(Activation('relu'))

    self.G.add(UpSampling2D())
    self.G.add(Conv2DTranspose(int(depth/4), 5, padding='same'))
    self.G.add(BatchNormalization(momentum=0.9))
    self.G.add(Activation('relu'))

    self.G.add(UpSampling2D())
    self.G.add(Conv2DTranspose(int(depth/8), 5, padding='same'))
    self.G.add(BatchNormalization(momentum=0.9))
    self.G.add(Activation('relu'))

    # Out: 64 x 64 x 1 grayscale image [0.0,1.0] per pix
    self.G.add(Conv2DTranspose(1, 5, padding='same'))
    self.G.add(Activation('sigmoid'))
    self.G.summary()
    return self.G
```

## Source

- https://

```
def discriminator_model(self):
    if self.DM:
        return self.DM
    optimizer = RMSprop(lr=2e-4, decay=1e-8)
    self.DM = Sequential()
    self.DM.add(self.discriminator())

    #multi_gpu
    with tf.device('/cpu:0'):
        self.DM = multi_gpu_model(self.DM, gpus=4)

    self.DM.compile(loss='binary_crossentropy', optimizer=optimizer, metrics=['accuracy'])
    return self.DM

def adversarial_model(self):
    if self.AM:
        return self.AM
    optimizer = RMSprop(lr=1e-4, decay=1e-8)
    self.AM = Sequential()
    self.AM.add(self.generator())
    self.AM.add(self.discriminator())

    #multi_gpu
    with tf.device('/cpu:0'):
        self.AM = multi_gpu_model(self.AM, gpus=4)

    self.AM.compile(loss='binary_crossentropy', optimizer=optimizer, metrics=['accuracy'])
    return self.AM
```

## Source

- <https://>

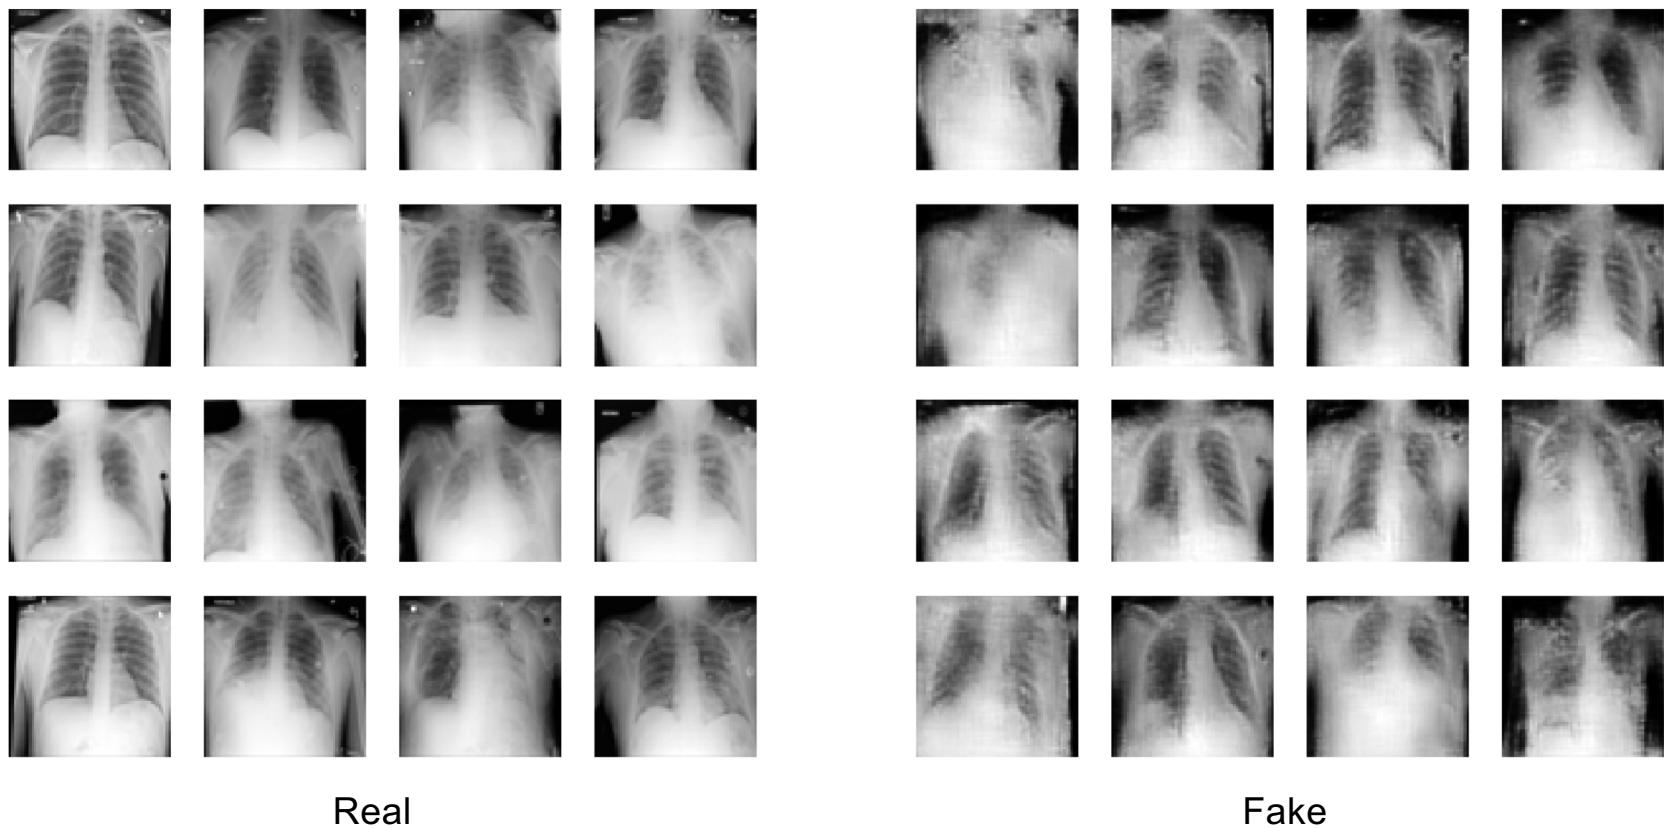
```
def train(self, train_steps=2000, batch_size=256, save_interval=0):
    print(self.x_train.shape[0])
    noise_input = None
    if save_interval>0:
        noise_input = np.random.uniform(-1.0, 1.0, size=[16, 100])
    for i in range(train_steps):
        images_train = self.x_train[np.random.randint(0,
            self.x_train.shape[0], size=batch_size), :, :, :]
        #train the discriminator first
        noise = np.random.uniform(-1.0, 1.0, size=[batch_size, 100])
        images_fake = self.generator.predict(noise)
        x = np.concatenate((images_train, images_fake))
        y = np.ones([2*batch_size, 1])
        y[batch_size:, :] = 0
        d_loss = self.discriminator.train_on_batch(x, y)

        #train the adversarial model second
        y = np.ones([batch_size, 1])
        noise = np.random.uniform(-1.0, 1.0, size=[batch_size, 100])
        a_loss = self.adversarial.train_on_batch(noise, y)

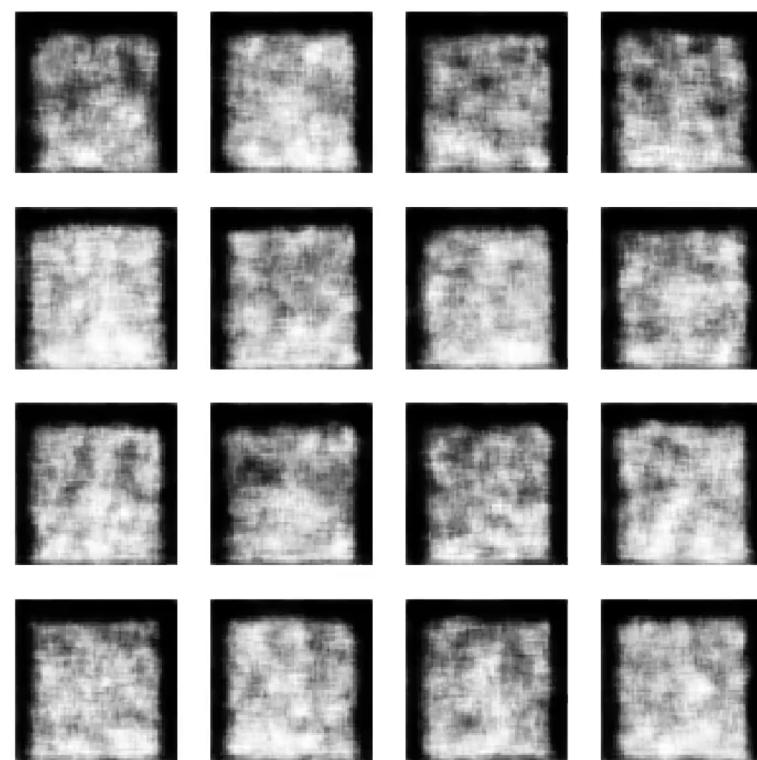
        log_mesg = "%d: [D loss: %f, acc: %f]" % (i, d_loss[0], d_loss[1])
        log_mesg = "%s [A loss: %f, acc: %f]" % (log_mesg, a_loss[0], a_loss[1])
        print(log_mesg)
        if save_interval>0:
            if (i+1)%save_interval==0:
                self.plot_images(save2file=True, samples=noise_input.shape[0],\
                    noise=noise_input, step=(i+1))
```

R8.py

## Results (2,126 samples, 2,000 steps, 18 minutes)



## Training Snapshots Every 10 Steps



## Test 2: GAN -> 128x128 synthetic images

- Training data
  - ‘No Finding’ subset
    - Male
    - Ages between 19 ~ 71
  - 19,011 samples
    - Down-sampled to 128x128 using imagemagick
- Training statistics
  - Four Nvidia Quadro M6000 GPUs
  - 2,000 steps, 256 samples / step
  - Finished in 36 minutes

## Source Code

- [https://github.com/abcsFrederick/gan/blob/master/scripts/NIH.AI\\_DCGAN\\_128x128\\_CXR8.py](https://github.com/abcsFrederick/gan/blob/master/scripts/NIH.AI_DCGAN_128x128_CXR8.py)

```
# (W-F+2P)/S+1
def discriminator(self):
    if self.D:
        return self.D
    self.D = Sequential()
    depth = 64
    dropout = 0.2

    input_shape = (self.img_rows, self.img_cols, self.channel)
    self.D.add(Conv2D(depth*1, 5, strides=2, input_shape=input_shape, padding='same'))
    self.D.add(LeakyReLU(alpha=0.2))
    self.D.add(Dropout(dropout))

    self.D.add(Conv2D(depth*2, 5, strides=2, padding='same'))
    self.D.add(LeakyReLU(alpha=0.2))
    self.D.add(Dropout(dropout))

    self.D.add(Conv2D(depth*4, 5, strides=2, padding='same'))
    self.D.add(LeakyReLU(alpha=0.2))
    self.D.add(Dropout(dropout))

    self.D.add(Conv2D(depth*2, 5, strides=1, padding='same'))
    self.D.add(LeakyReLU(alpha=0.2))
    self.D.add(Dropout(dropout))

    # Out: 1-dim probability
    self.D.add(Flatten())
    self.D.add(Dense(1))
    self.D.add(Activation('sigmoid'))
    self.D.summary()
    return self.D
```

## Source Code

- <https://github.com>

```
def generator(self):
    if self.G:
        return self.G
    self.G = Sequential()
    dropout = 0.2
    depth = 64+64+64+64
    dim = 8
    # In: 100
    # Out: dim x dim x depth
    self.G.add(Dense(dim*dim*depth, input_dim=100))
    self.G.add(BatchNormalization(momentum=0.9))
    self.G.add(Activation('relu'))
    self.G.add(Reshape((dim, dim, depth)))
    self.G.add(Dropout(dropout))

    self.G.add(UpSampling2D())
    self.G.add(Conv2DTranspose(int(depth), 5, padding='same'))
    self.G.add(BatchNormalization(momentum=0.9))
    self.G.add(Activation('relu'))

    self.G.add(UpSampling2D())
    self.G.add(Conv2DTranspose(int(depth/2), 5, padding='same'))
    self.G.add(BatchNormalization(momentum=0.9))
    self.G.add(Activation('relu'))

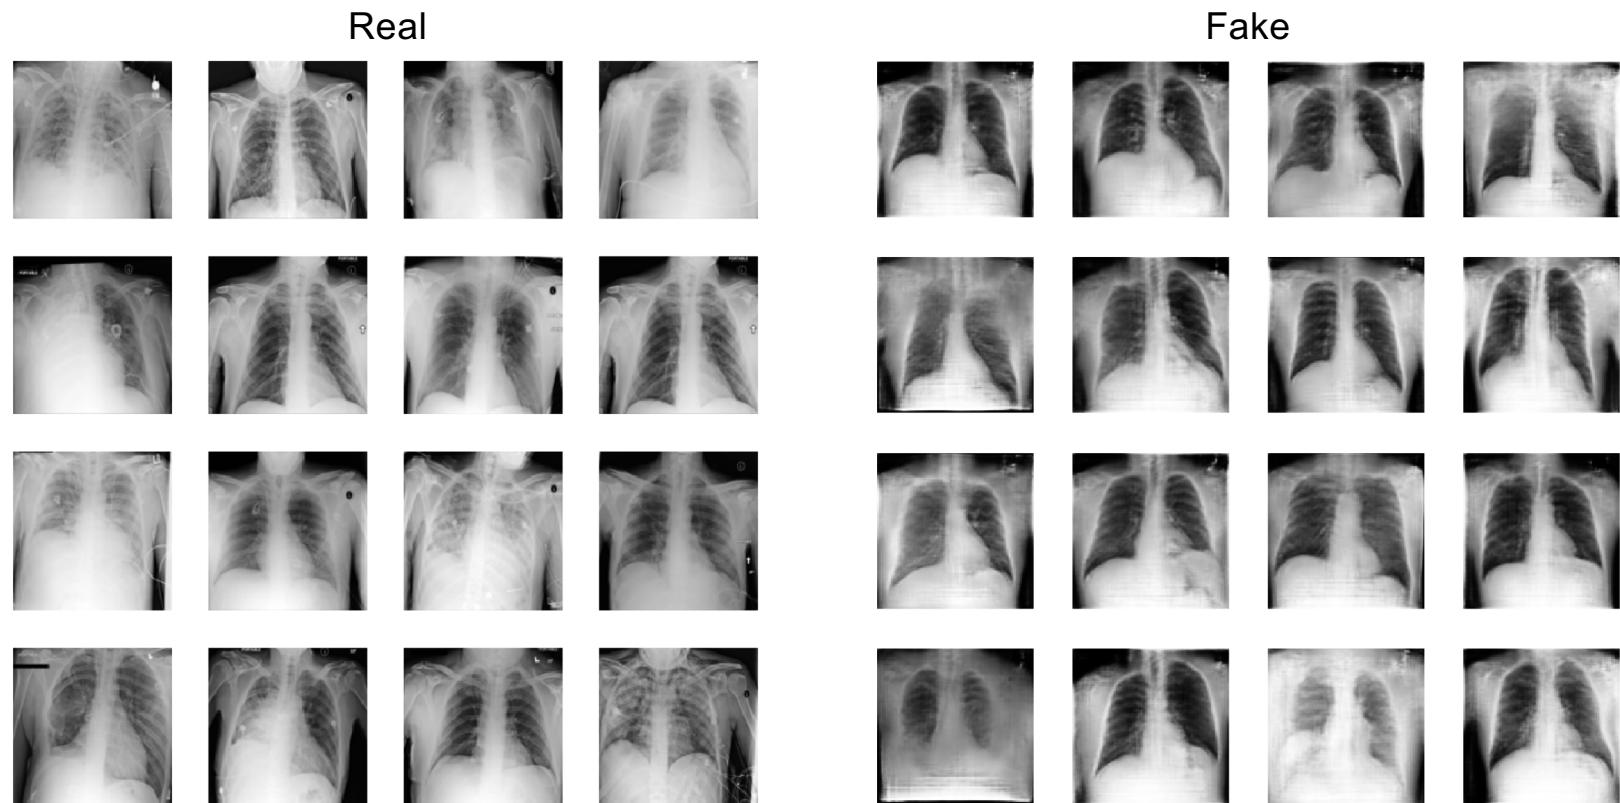
    self.G.add(UpSampling2D())
    self.G.add(Conv2DTranspose(int(depth/4), 5, padding='same'))
    self.G.add(BatchNormalization(momentum=0.9))
    self.G.add(Activation('relu'))

    self.G.add(UpSampling2D())
    self.G.add(Conv2DTranspose(int(depth/8), 5, padding='same'))
    self.G.add(BatchNormalization(momentum=0.9))
    self.G.add(Activation('relu'))

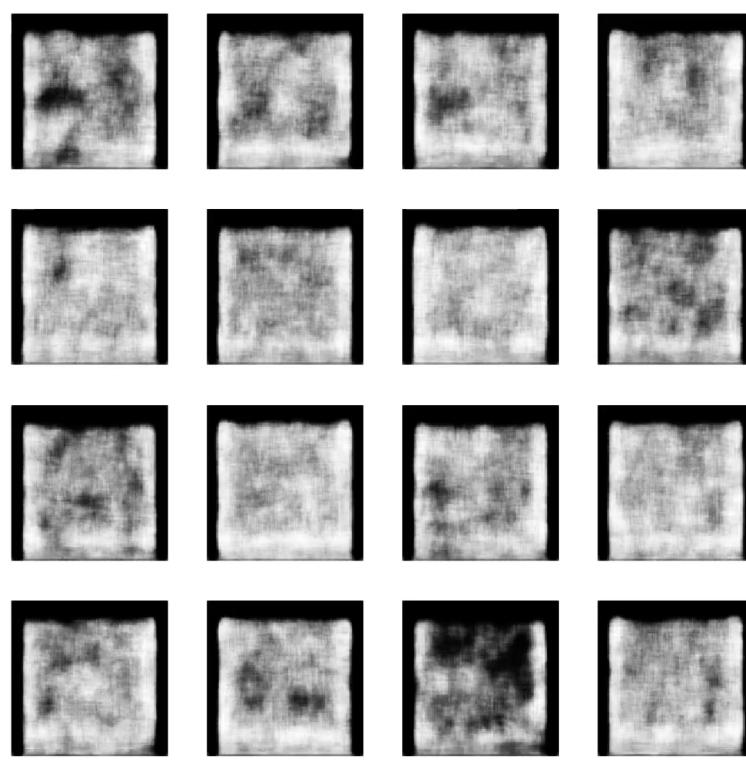
    # Out: 128 x 128 x 1 grayscale image [0.0,1.0] per pix
    self.G.add(Conv2DTranspose(1, 5, padding='same'))
    self.G.add(Activation('sigmoid'))
    self.G.summary()
    return self.G
```

XR8.py

## Results (19,011 samples, 2,000 steps, 36 minutes)



## Training Snapshots Every 10 Steps



## Conclusion

64x64



128x128

